

Project Phase 2

CSC 555 Mining Big Data
Michal Chowaniak

(due Sunday June 16th)

In this part of the project, you will execute queries using Hive, Pig and Hadoop streaming and develop a custom version of KMeans clustering. The schema is available below, but don't forget to change to the correct delimiter:

http://rasinsrv07.cstcis.cti.depaul.edu/CSC555/SSBM1/SSBM_schema_hive.sql

The data is available at (this is Scale1, the smallest denomination of this benchmark)

<http://rasinsrv07.cstcis.cti.depaul.edu/CSC555/SSBM1/>

In your submission, please note what cluster you are using. Please be sure to submit all code (pig, python and Hive). You should also submit the command lines you use and a screenshot of a completed run (just the last page, do not worry about capturing the whole output). An answer without corresponding code will not be counted.

I highly recommend creating a small sample input (e.g., by running `head lineorder.tbl > lineorder.tbl.sample`, you can create a small version of lineorder with a few lines) and testing your code with it. You can run `head -n 500 lineorder.tbl` to get a specific number of lines.

NOTE: the total number of points adds up to 70.

At the begging I used 4 node cluster, but after I turned off 2 instances, I was not able to get them back, so some of answers were run on 2 node cluster.

In operation

| Node | Last contact | Admin State | Capacity | Used | Non DFS Used | Remaining | Blocks | Block pool used | Failed Volumes | Version |
|---|--------------|-------------|----------|---------|--------------|-----------|--------|-----------------|----------------|---------|
| ip-172-31-3-227.us-east-2.compute.internal (172.31.3.227:50010) | 0 | In Service | 29.99 GB | 2.18 GB | 2.28 GB | 25.52 GB | 286 | 2.18 GB (7.27%) | 0 | 2.6.4 |
| ip-172-31-12-205.us-east-2.compute.internal (172.31.12.205:50010) | 1 | In Service | 29.99 GB | 2.18 GB | 5.81 GB | 22 GB | 286 | 2.18 GB (7.27%) | 0 | 2.6.4 |

Part 1: Data Transformation (15 pts)

Transform part.tbl table into a ~-separated ('~') file: Use Hive, MapReduce with HadoopStreaming and Pig (i.e. 3 different solutions).

In all solutions you must switch the first two columns (i.e., switch the positions of columns 1 and 2). You do not need to transform the columns in any way, just switch them around. Note that this means you do not have to use SELECT TRANSFORM or python in your Hive solution.

wget <http://rasinsrv07.cstcis.cti.depaul.edu/CSC555/SSBM1/part.tbl>

Hive

nano part.tbl

```
GNU nano 2.9.8 part.tbl
1|lace spring|MFGR#1|MFGR#11|MFGR#1121|goldenrod|PROMO BURNISHED COPPER|7|JUMBO PKG|
2|rosy metallic|MFGR#4|MFGR#43|MFGR#4318|blush|LARGE BRUSHED BRASS|1|LG CASE|
3|green antique|MFGR#3|MFGR#32|MFGR#3210|dark|STANDARD POLISHED BRASS|21|WRAP CASE|
```

nano SSBM_schema_hive.sql

```
create table part (
  p_partkey int,
  p_name varchar(22),
  p_mfgr varchar(6),
  p_category varchar(7),
  p_brand1 varchar(9),
  p_color varchar(11),
  p_type varchar(25),
  p_size int,
  p_container varchar(10))
ROW FORMAT DELIMITED FIELDS
TERMINATED BY '|' STORED AS TEXTFILE;
```

```

create table part (
  p_partkey      int,
  p_name         varchar(22),
  p_mfgr         varchar(6),
  p_category     varchar(7),
  p_brand1       varchar(9),
  p_color        varchar(11),
  p_type         varchar(25),
  p_size         int,
  p_container    varchar(10))
ROW FORMAT DELIMITED FIELDS
TERMINATED BY '|' STORED AS TEXTFILE;

```

LOAD DATA

```

LOAD DATA LOCAL INPATH '/home/ec2-user/part.tbl'
OVERWRITE INTO TABLE part;

```

```

> LOAD DATA LOCAL INPATH '/home/ec2-user/part.tbl'
> OVERWRITE INTO TABLE part;
Loading data to table default.part
OK
Time taken: 2.084 seconds
hive>

```

```

SELECT COUNT(*) FROM part;

```

```

Total MapReduce CPU Time Spent: 2 seconds 440 msec
OK
200000
Time taken: 16.094 seconds, Fetched: 1 row(s)
hive>

```

```

CREATE TABLE PartSwap10 (p_name String, p_partkey INT, p_mfgr String, p_category String, p_brand1 String,
p_color String, p_type String, p_size INT, p_container String)
ROW FORMAT DELIMITED FIELDS

```

TERMINATED BY '~' STORED AS TEXTFILE;

```
hive> CREATE TABLE PartSwap10 (p_name String, p_partkey INT, p_mfgr String, p_category String, p_brand1 String, p_color String, p_type String, p_size INT, p_container String)
> ROW FORMAT DELIMITED FIELDS
> TERMINATED BY '~' STORED AS TEXTFILE;
OK
Time taken: 0.046 seconds
hive>
```

INSERT OVERWRITE TABLE PartSwap10

SELECT p_name, p_partkey, p_mfgr, p_category, p_brand1,
p_color, p_type, p_size, p_container FROM part;

```
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Cumulative CPU: 3.01 sec HDFS Read: 17145000 HDFS Write: 16939343
SUCCESS
Total MapReduce CPU Time Spent: 3 seconds 10 msec
OK
Time taken: 12.048 seconds
hive>
```

hadoop fs -cat /user/hive/warehouse/partswap10/000000_0

```
aquamarine chiffon~199994~MFGR#1~MFGR#14~MFGR#1414~frosted~LARGE BRUSHED COPPER~32~MED CAN
dodger magenta~199995~MFGR#4~MFGR#45~MFGR#458~blanched~PROMO POLISHED TIN~50~WRAP CAN
steel cyan~199996~MFGR#4~MFGR#44~MFGR#4418~chocolate~PROMO PLATED COPPER~11~MED PACK
azure snow~199997~MFGR#4~MFGR#44~MFGR#4426~drab~PROMO PLATED NICKEL~37~SM DRUM
misty plum~199998~MFGR#5~MFGR#55~MFGR#5512~peach~MEDIUM BURNISHED BRASS~49~LG BOX
azure cream~199999~MFGR#5~MFGR#52~MFGR#5235~medium~PROMO PLATED BRASS~24~LG CASE
light midnight~200000~MFGR#5~MFGR#52~MFGR#5223~firebrick~MEDIUM ANODIZED TIN~22~LG CAN
[ec2-user@ip-172-31-12-205 ~]$
```

MapReduce with HadoopStreaming

```
ec2-user@ip-172-31-12-205:~  
GNU nano 2.9.8 part1Mapper.py  
#!/usr/bin/python  
import sys  
  
reorder = []  
  
for line in sys.stdin:  
    line = line.strip()  
    #split = line.split('|')  
    split = line.split('~')  
    reorder = split[1],split[0],split[2],split[3],split[4],split[5],split[6],split[7],split[8]  
    # print '\t'.join(reorder)  
    print '~'.join(reorder)
```

cat part.tbl | python part1Mapper.py

```
steel cyan~199996~MFGR#4~MFGR#44~MFGR#4418~chocolate~PROMO PLATED COPPER~11~  
azure snow~199997~MFGR#4~MFGR#44~MFGR#4426~drab~PROMO PLATED NICKEL~37~  
misty plum~199998~MFGR#5~MFGR#55~MFGR#5512~peach~MEDIUM BURNISHED BRASS~49~  
azure cream~199999~MFGR#5~MFGR#52~MFGR#5235~medium~PROMO PLATED BRASS~24~  
light midnight~200000~MFGR#5~MFGR#52~MFGR#5223~firebrick~MEDIUM ANODIZED TIN~22~  
[ec2-user@ip-172-31-12-205 ~]$
```

hadoop jar /home/ec2-user/hadoop-2.6.4/hadoop-streaming-2.6.4.jar -input /user/ec2-user/part.tbl -output /user/ec2-user/partreorder02 -mapper part1Mapper.py -file /home/ec2-user/part1Mapper.py

```
IO_ERROR=0  
WRONG_LENGTH=0  
WRONG_MAP=0  
WRONG_REDUCE=0  
File Input Format Counters  
    Bytes Read=17143355  
File Output Format Counters  
    Bytes Written=17139259  
19/06/12 05:13:03 INFO streaming.StreamJob: Output directory: /user/ec2-user/partreorder02  
[ec2-user@ip-172-31-12-205 ~]$
```

hadoop fs -cat /user/ec2-user/partreorder02/part-00000

```
yellow white~3722~MFGR#4~MFGR#41~MFGR#4110~honeydew~STANDARD PLATED NICKEL~38~SM BOX  
yellow white~50230~MFGR#3~MFGR#35~MFGR#359~seashell~MEDIUM POLISHED NICKEL~14~MED BOX  
yellow white~56303~MFGR#5~MFGR#52~MFGR#5220~tan~ECONOMY ANODIZED STEEL~28~WRAP PKG  
yellow white~90915~MFGR#3~MFGR#33~MFGR#3319~red~SMALL PLATED BRASS~6~MED JAR  
[ec2-user@ip-172-31-12-205 ~]$
```

Pig

```
cd $PIG_HOME
```

```
bin/pig
```

```
hadoop fs -put part.tbl /user/ec2-user/
```

```
[ec2-user@ip-172-31-12-205 ~]$ hadoop fs -ls /user/ec2-user/
Found 6 items
drwxr-xr-x  - ec2-user supergroup          0 2019-06-08 07:05 /user/ec2-user/als
drwxr-xr-x  - ec2-user supergroup          0 2019-06-08 06:38 /user/ec2-user/dataset
drwxr-xr-x  - ec2-user supergroup          0 2019-06-08 06:39 /user/ec2-user/ml_dataset
drwxr-xr-x  - ec2-user supergroup          0 2019-06-08 06:35 /user/ec2-user/movielens
-rw-r--r--  2 ec2-user supergroup 17139259 2019-06-11 02:36 /user/ec2-user/part.tbl
drwxr-xr-x  - ec2-user supergroup          0 2019-06-08 07:09 /user/ec2-user/recommendatio
ns
[ec2-user@ip-172-31-12-205 ~]$
```

```
PartData = LOAD '/user/ec2-user/part.tbl' USING PigStorage('|')
AS (p_partkey:INT, p_name:CHARARRAY, p_mfgr:CHARARRAY, p_category:CHARARRAY,
p_brand1:CHARARRAY, p_color:CHARARRAY, p_type:CHARARRAY, p_size:INT, p_container:CHARARRAY);
```

```
DESCRIBE PartData;
```

```
grunt> PartData = LOAD '/user/ec2-user/part.tbl' USING PigStorage('|')
>> AS (p_partkey:INT, p_name:CHARARRAY, p_mfgr:CHARARRAY, p_category:CHARARRAY, p_brand1:CH
ARARRAY, p_color:CHARARRAY, p_type:CHARARRAY, p_size:INT, p_container:CHARARRAY);
2019-06-11 03:13:30,822 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.
default.name is deprecated. Instead, use fs.defaultFS
grunt> DESCRIBE PartData;
PartData: {p_partkey: int,p_name: chararray,p_mfgr: chararray,p_category: chararray,p_brand
1: chararray,p_color: chararray,p_type: chararray,p_size: int,p_container: chararray}
grunt>
```

```
PartDataSwitched = FOREACH PartData GENERATE p_name, p_partkey, p_mfgr, p_category, p_brand1,
p_color, p_type, p_size, p_container;
```

```
DUMP PartDataSwitched;
```

```
(dodger magenta,199995,MFGR#4,MFGR#45,MFGR#458,blanched,PROMO POLISHED TIN,50,WRAP CAN)
(steel cyan,199996,MFGR#4,MFGR#44,MFGR#4418,chocolate,PROMO PLATED COPPER,11,MED PACK)
(azure snow,199997,MFGR#4,MFGR#44,MFGR#4426,drab,PROMO PLATED NICKEL,37,SM DRUM)
(misty plum,199998,MFGR#5,MFGR#55,MFGR#5512,peach,MEDIUM BURNISHED BRASS,49,LG BOX)
(azure cream,199999,MFGR#5,MFGR#52,MFGR#5235,medium,PROMO PLATED BRASS,24,LG CASE)
(light midnight,200000,MFGR#5,MFGR#52,MFGR#5223,firebrick,MEDIUM ANODIZED TIN,22,LG CAN)
grunt>
```

STORE PartDataSwitched INTO 'PartDataNewSwitched2' USING PigStorage('~');

```
SourceManager at 172.31.12.205:8032
2019-06-12 05:21:34,834 [main] INFO org.apache.hadoop.mapred.ClientServiceDelegate - Applica
tion state is completed. FinalApplicationStatus=SUCCEEDED. Redirecting to job history server
2019-06-12 05:21:34,861 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceL
ayer.MapReduceLauncher - Success!
grunt> █
```

hadoop fs -cat /user/ec2-user/PartDataNewSwitched2/part-m-00000

```
aquamarine chiffon~199994~MFGR#1~MFGR#14~MFGR#1414~frosted~LARGE BRUSHED COPPER~32~MED CAN
dodger magenta~199995~MFGR#4~MFGR#45~MFGR#458~blanched~PROMO POLISHED TIN~50~WRAP CAN
steel cyan~199996~MFGR#4~MFGR#44~MFGR#4418~chocolate~PROMO PLATED COPPER~11~MED PACK
azure snow~199997~MFGR#4~MFGR#44~MFGR#4426~drab~PROMO PLATED NICKEL~37~SM DRUM
misty plum~199998~MFGR#5~MFGR#55~MFGR#5512~peach~MEDIUM BURNISHED BRASS~49~LG BOX
azure cream~199999~MFGR#5~MFGR#52~MFGR#5235~medium~PROMO PLATED BRASS~24~LG CASE
light midnight~200000~MFGR#5~MFGR#52~MFGR#5223~firebrick~MEDIUM ANODIZED TIN~22~LG CAN
[ec2-user@ip-172-31-12-205 pig-0.15.0]$ █
```

Part 2: Querying (25 pts)

Implement the following query:

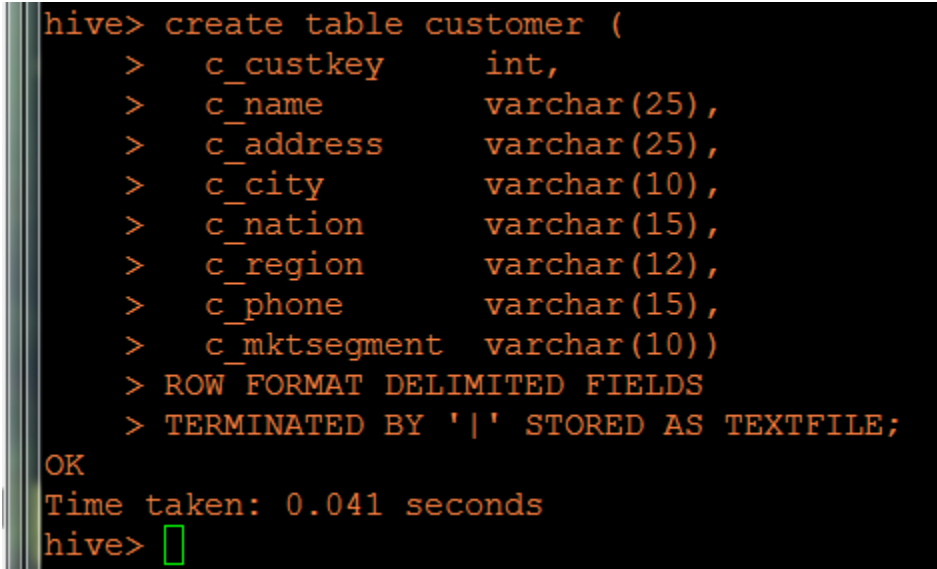
```
select c_nation, sum(lo_revenue)
from customer, lineorder
where lo_custkey = c_custkey
and c_region = 'AMERICA'
and lo_discount BETWEEN 4 and 6
group by c_nation;
```

using Hive, MapReduce with HadoopStreaming and Pig (i.e. 3 different solutions). I Hive, this merely requires pasting the query into the Hive prompt and timing it. In Hadoop streaming, this will require a total of 2 passes (one for join and another one for GROUP BY).

Hive

```
wget http://rasinsrv07.cstcis.cti.depaul.edu/CSC555/SSBM1/customer.tbl
wget http://rasinsrv07.cstcis.cti.depaul.edu/CSC555/SSBM1/lineorder.tbl
```

```
create table customer (
  c_custkey  int,
  c_name     varchar(25),
  c_address  varchar(25),
  c_city     varchar(10),
  c_nation   varchar(15),
  c_region   varchar(12),
  c_phone    varchar(15),
  c_mktsegment varchar(10))
ROW FORMAT DELIMITED FIELDS
TERMINATED BY '|' STORED AS TEXTFILE;
```



```
hive> create table customer (
>   c_custkey      int,
>   c_name         varchar(25),
>   c_address      varchar(25),
>   c_city         varchar(10),
>   c_nation       varchar(15),
>   c_region       varchar(12),
>   c_phone        varchar(15),
>   c_mktsegment   varchar(10))
> ROW FORMAT DELIMITED FIELDS
> TERMINATED BY '|' STORED AS TEXTFILE;
OK
Time taken: 0.041 seconds
hive> 
```

```
create table lineorder (
  lo_orderkey  int,
  lo_linenumbr int,
  lo_custkey   int,
  lo_partkey   int,
  lo_suppkey   int,
  lo_orderdate int,
  lo_orderpriority varchar(15),
```



```
lo_shippriority   varchar(1),
lo_quantity       int,
lo_extendedprice  int,
lo_ordertotalprice int,
lo_discount       int,
lo_revenue        int,
lo_supplycost     int,
lo_tax            int,
lo_commitdate     int,
lo_shipmode       varchar(10))
ROW FORMAT DELIMITED FIELDS
TERMINATED BY '|' STORED AS TEXTFILE;
```

```
hive> create table lineorder (
  >   lo_orderkey      int,
  >   lo_linenumbers   int,
  >   lo_custkey       int,
  >   lo_partkey       int,
  >   lo_suppkey       int,
  >   lo_orderdate     int,
  >   lo_orderpriority varchar(15),
  >   lo_shippriority  varchar(1),
  >   lo_quantity      int,
  >   lo_extendedprice int,
  >   lo_ordertotalprice int,
  >   lo_discount      int,
  >   lo_revenue       int,
  >   lo_supplycost    int,
  >   lo_tax           int,
  >   lo_commitdate    int,
  >   lo_shipmode      varchar(10))
  > ROW FORMAT DELIMITED FIELDS
  > TERMINATED BY '|' STORED AS TEXTFILE;
OK
Time taken: 0.041 seconds
hive> █
```

LOAD DATA LOCAL INPATH '/home/ec2-user/lineorder.tbl' OVERWRITE INTO TABLE lineorder;

```
hive> LOAD DATA LOCAL INPATH '/home/ec2-user/lineorder.tbl' OVERWRITE INTO TABLE lineorder;
Loading data to table default.lineorder
OK
Time taken: 9.017 seconds
hive>
```

LOAD DATA LOCAL INPATH '/home/ec2-user/customer.tbl' OVERWRITE INTO TABLE customer;

```
hive> LOAD DATA LOCAL INPATH '/home/ec2-user/customer.tbl' OVERWRITE INTO TABLE customer;
Loading data to table default.customer
OK
Time taken: 0.133 seconds
hive>
```

```
select c_nation, sum(lo_revenue)
from customer, lineorder
where lo_custkey = c_custkey
and c_region = 'AMERICA'
and lo_discount BETWEEN 4 and 6
group by c_nation;
```

```
Stage-Stage-2: Map: 3 Reduce: 3 Cumulative CPU: 21.66 s
ite: 108 SUCCESS
Total MapReduce CPU Time Spent: 21 seconds 660 msec
OK
ARGENTINA      243988697072
BRAZIL  225595365795
UNITED STATES  244263170830
CANADA  240715548308
PERU      228441124985
Time taken: 29.564 seconds, Fetched: 5 row(s)
hive>
```

MapReduce with HadoopStreaming

```
select c_nation, sum(lo_revenue)
from customer, lineorder
where lo_custkey = c_custkey
and c_region = 'AMERICA'
and lo_discount BETWEEN 4 and 6
group by c_nation;
```

wget <http://rasinsrv07.cstcis.cti.depaul.edu/CSC555/SSBM1/customer.tbl>

wget <http://rasinsrv07.cstcis.cti.depaul.edu/CSC555/SSBM1/lineorder.tbl>

head customer.tbl

```
[ec2-user@ip-172-31-12-205 ~]$ head customer.tbl
1|Customer#000000001|j5JsirBM9P|MOROCCO 0|MOROCCO|AFRICA|25-989-741-2988|BUILDING|
2|Customer#000000002|487LWldovn6Q4dMVym|JORDAN 1|JORDAN|MIDDLE EAST|23-768-687-3665|AUTOMOBILE|
3|Customer#000000003|fkRGN8n|ARGENTINA7|ARGENTINA|AMERICA|11-719-748-3364|AUTOMOBILE|
4|Customer#000000004|4u58h f|EGYPT 4|EGYPT|MIDDLE EAST|14-128-190-5944|MACHINERY|
5|Customer#000000005|hwBtxkoBF qSW4KrI|CANADA 5|CANADA|AMERICA|13-750-942-6364|HOUSEHOLD|
6|Customer#000000006| g1s,pzDenUEBW3O,2 pxu|SAUDI ARA2|SAUDI ARABIA|MIDDLE EAST|30-114-968-4951|AUTOMOBILE|
7|Customer#000000007|8OkMVLQldK6Mbu6WG9|CHINA 0|CHINA|ASIA|28-190-982-9759|AUTOMOBILE|
8|Customer#000000008|j,pZ,Qp,qtFEo0r0c 92qo|PERU 6|PERU|AMERICA|27-147-574-9335|BUILDING|
9|Customer#000000009|vgIql8H6zoyuLMFN|INDIA 6|INDIA|ASIA|18-338-906-3675|FURNITURE|
10|Customer#000000010|Vf mQ6Ug9Ucf5OKGYq fs|ETHIOPIA 9|ETHIOPIA|AFRICA|15-741-346-9870|HOUSEHOLD|
[ec2-user@ip-172-31-12-205 ~]$
```

```
create table customer (
  c_custkey      int,
  c_name         varchar(25),
  c_address      varchar(25),
  c_city         varchar(10),
  c_nation       varchar(15),
  c_region       varchar(12),
  c_phone        varchar(15),
  c_mktsegment   varchar(10)
);
```

head lineorder.tbl

```
[ec2-user@ip-172-31-12-205 ~]$ head lineorder.tbl
1|1|7381|155190|828|19960102|5-LOW|0|17|2116823|17366547|4|2032150|74711|2|19960212|TRUCK|
1|2|7381|67310|163|19960102|5-LOW|0|36|4598316|17366547|9|4184467|76638|6|19960228|MAIL|
1|3|7381|63700|71|19960102|5-LOW|0|8|1330960|17366547|10|1197864|99822|2|19960305|REG AIR|
1|4|7381|2132|943|19960102|5-LOW|0|28|2895564|17366547|9|2634963|62047|6|19960330|AIR|
1|5|7381|24027|1625|19960102|5-LOW|0|24|2282448|17366547|10|2054203|57061|4|19960314|FOB|
1|6|7381|15635|1368|19960102|5-LOW|0|32|4962016|17366547|7|4614674|93037|2|19960207|MAIL|
2|1|15601|106170|1066|19961201|1-URGENT|0|38|4469446|4692918|0|4469446|70570|5|19970114|RAIL|
3|1|24664|4297|1959|19931014|5-LOW|0|45|5405805|19384625|6|5081456|72077|0|19940104|AIR|
3|2|24664|19036|1667|19931014|5-LOW|0|49|4679647|19384625|10|4211682|57301|0|19931220|RAIL|
3|3|24664|128449|1409|19931014|5-LOW|0|27|3989088|19384625|6|3749742|88646|7|19931122|SHIP|
[ec2-user@ip-172-31-12-205 ~]$
```

```
create table lineorder (
  lo_orderkey      int,
  lo_linenumbers   int,
  lo_custkey       int,
  lo_partkey       int,
  lo_suppkey       int,
  lo_orderdate     int,
  lo_orderpriority varchar(15),
  lo_shippriority  varchar(1),
  lo_quantity      int,
  lo_extendedprice int,
  lo_ordertotalprice int,
  lo_discount      int,
  lo_revenue       int,
  lo_supplycost    int,
  lo_tax           int,
  lo_commitdate    int,
  lo_shipmode      varchar(10)
);
```

hadoop fs -mkdir -p /data/joinCustomer

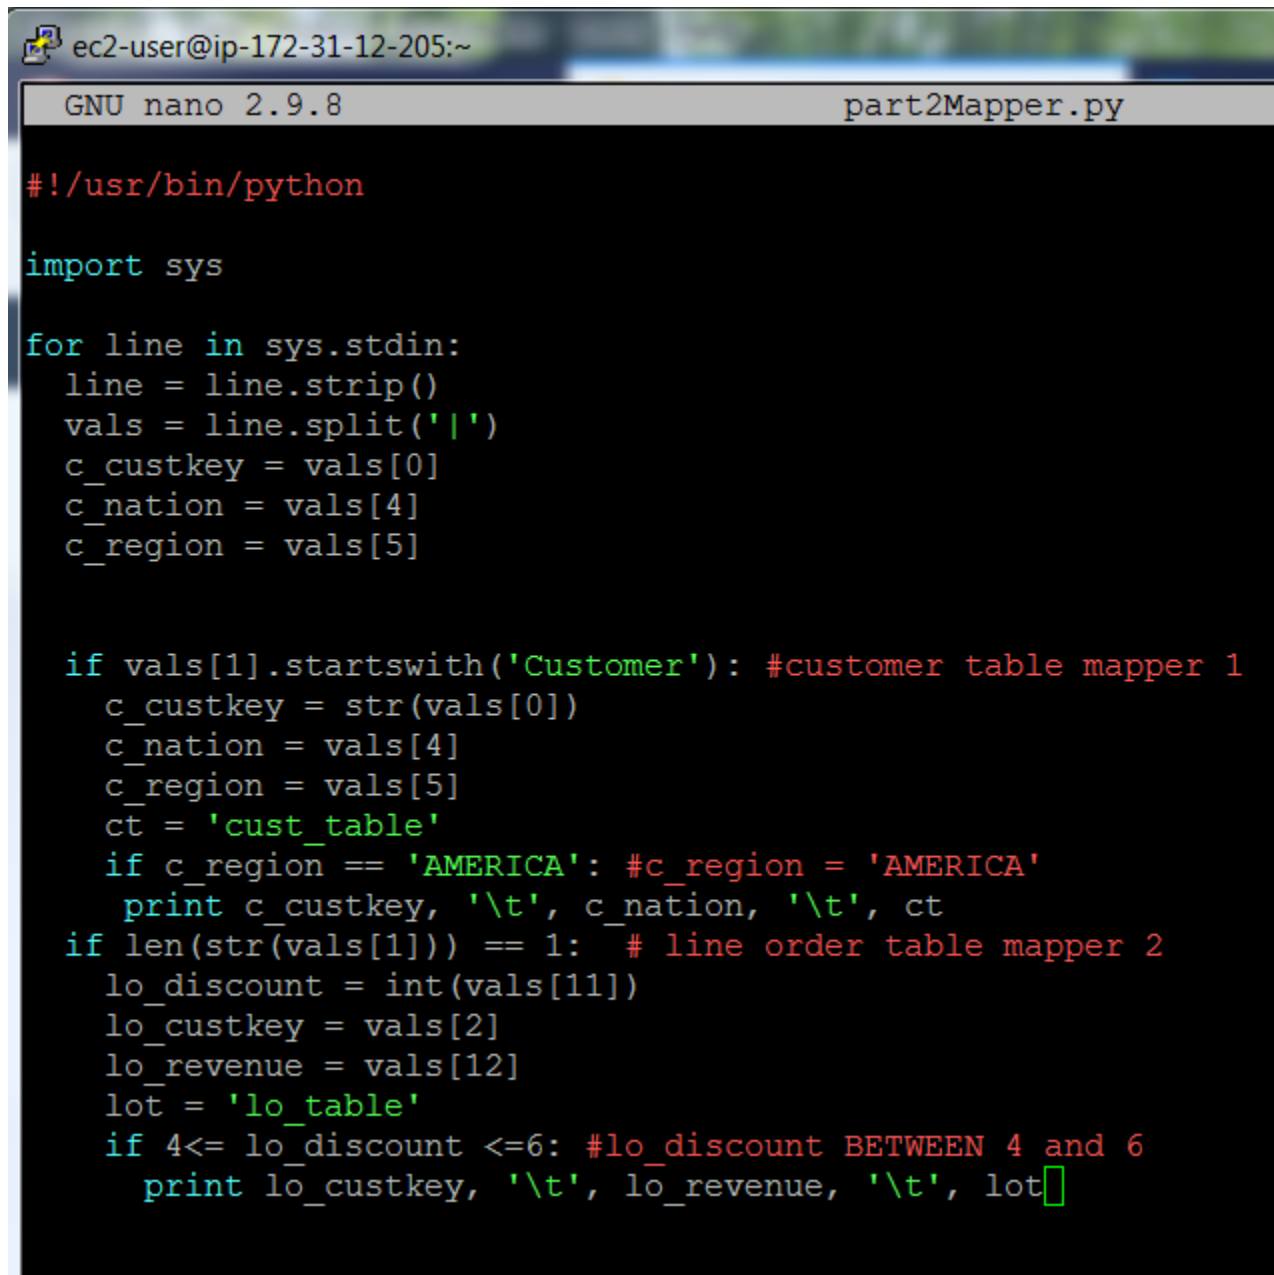
```
[ec2-user@ip-172-31-12-205 ~]$ hadoop fs -ls /data/ | grep join
drwxr-xr-x  - ec2-user supergroup          0 2019-06-15 03:47 /data/joinCustomer
[ec2-user@ip-172-31-12-205 ~]$
```

hadoop fs -put customer.tbl lineorder.tbl /data/joinCustomer

```
[ec2-user@ip-172-31-12-205 ~]$ hadoop fs -ls /data/joinCustomer/
Found 2 items
-rw-r--r--  2 ec2-user supergroup    2837046 2019-06-15 03:49 /data/joinCustomer/customer.
tbl
-rw-r--r--  2 ec2-user supergroup    594313001 2019-06-15 03:49 /data/joinCustomer/lineorder
.tbl
[ec2-user@ip-172-31-12-205 ~]$
```

Pass 1

part2Mapper.py



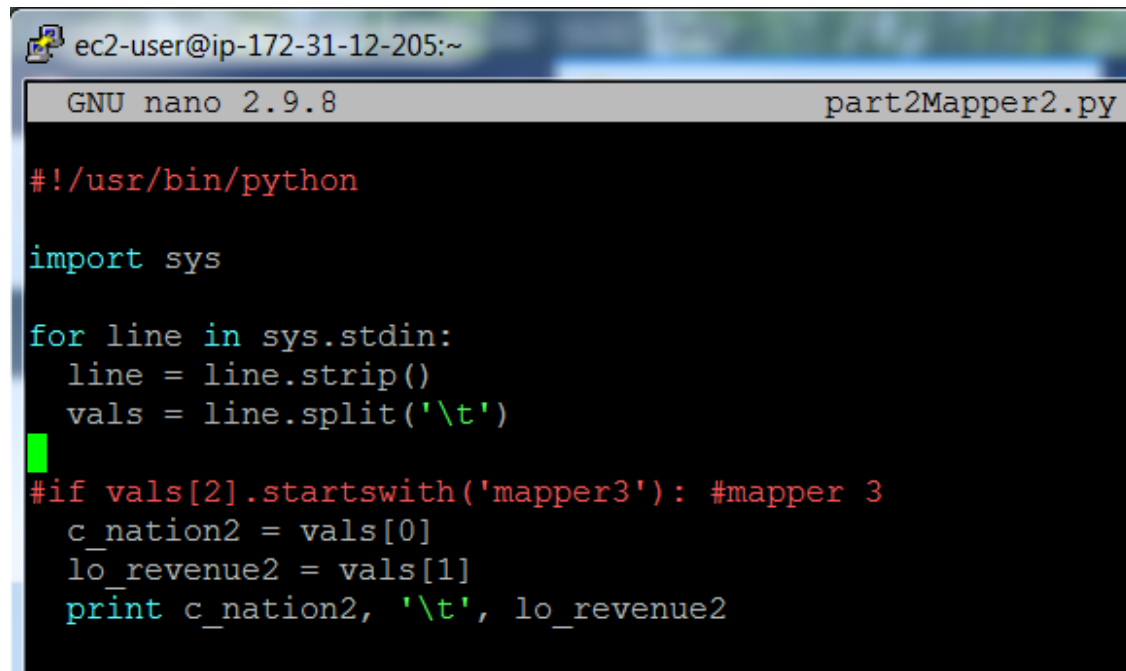
```
ec2-user@ip-172-31-12-205:~  
GNU nano 2.9.8 part2Mapper.py  
#!/usr/bin/python  
  
import sys  
  
for line in sys.stdin:  
    line = line.strip()  
    vals = line.split('|')  
    c_custkey = vals[0]  
    c_nation = vals[4]  
    c_region = vals[5]  
  
    if vals[1].startswith('Customer'): #customer table mapper 1  
        c_custkey = str(vals[0])  
        c_nation = vals[4]  
        c_region = vals[5]  
        ct = 'cust_table'  
        if c_region == 'AMERICA': #c_region = 'AMERICA'  
            print c_custkey, '\t', c_nation, '\t', ct  
    if len(str(vals[1])) == 1: # line order table mapper 2  
        lo_discount = int(vals[11])  
        lo_custkey = vals[2]  
        lo_revenue = vals[12]  
        lot = 'lo_table'  
        if 4<= lo_discount <=6: #lo_discount BETWEEN 4 and 6  
            print lo_custkey, '\t', lo_revenue, '\t', lot
```

part2Reducer.py

```
ec2-user@ip-172-31-12-205:~  
GNU nano 2.9.8 part2Reducer.py  
  
#!/usr/bin/python  
import sys  
curr_id = None  
curr_cnt = 0  
id = None  
lo_table_vals = None  
cust_table_vals = None  
custkey=[]  
lokey=[]  
rev_sub_sum = float(0)  
  
for line in sys.stdin:  
    line = line.strip()  
    ln = line.split('\t')  
    id = int(ln[0])  
    if curr_id == id: # key is the same  
        if ln[2].startswith('lo_table'): #checks what reducers get and puts to appropia  
            lokey.append(id) # to do inner join  
            lo_rev = float(ln[1]) # extract revenue value from 2nd mapper output  
            rev_sub_sum = float(rev_sub_sum)  
            rev_sub_sum = rev_sub_sum + lo_rev #sum revenues when while key is the same  
        if ln[2].startswith('cust_table'):  
            c_nation = ln[1]  
            custkey.append(id) # to do inner join  
    else: #new key has arrived  
        if curr_id: # output the count, single key completed  
            lolen = len(lokey) #check the lenght of list  
            custlen = len(custkey)  
            if (lolen * custlen > 0): # inner join  
                rev_sub_sum = str(int(rev_sub_sum))  
                print c_nation, '\t', rev_sub_sum, '\t', 'mapper3'  
            curr_id = id # update key  
        if ln[2].startswith('lo_table'):  
            lokey = []  
            lokey.append(ln[1])  
            custkey = []  
        if ln[2].startswith('cust_table'):  
            lokey = []  
            custkey = []  
            custkey.append(ln[1])  
lolenlast = len(lokey)  
custlenlast = len(custkey)  
if (lolenlast * custlenlast > 0): # inner join  
    rev_sub_sum = str(int(rev_sub_sum))  
    print c_nation, '\t', rev_sub_sum, '\t', 'mapper3'
```

Pass 2

part2Mapper2.py



```
ec2-user@ip-172-31-12-205:~  
GNU nano 2.9.8 part2Mapper2.py  
#!/usr/bin/python  
  
import sys  
  
for line in sys.stdin:  
    line = line.strip()  
    vals = line.split('\t')  
  
    #if vals[2].startswith('mapper3'): #mapper 3  
        c_nation2 = vals[0]  
        lo_revenue2 = vals[1]  
        print c_nation2, '\t', lo_revenue2
```

part2Reducer2.py

part2Reducer2.py

```
ec2-user@ip-172-31-12-205:~  
GNU nano 2.9.8 part2Reducer2.py  
  
#!/usr/bin/python  
import sys  
  
curr_id = None  
id = None  
curr_sum = 0  
  
# The input comes from standard input (line by line)  
for line in sys.stdin:  
    line = line.strip()  
    # parse the line and split it by '\t'  
    ln = line.split('\t')  
    # grab the key (int)  
    id = ln[0] # extract nation from print line  
    revenue = int(ln[1]) # extract revenue from print line  
    if curr_id == id: # key is the same  
        curr_sum += revenue  
    else: #new key has arrived  
        if curr_id: # output the count, single key completed  
            print curr_id, '\t', curr_sum  
        curr_id = id # update key  
        curr_sum = 0  
  
if curr_id == id:  
    curr_sum += revenue  
    print curr_id, '\t', curr_sum
```


All 4 python files (part2Mapper.py, part2Reducer.py, part2Mapper2.py, part2Reducer2.py) run without errors in command line using below command.

cat customer.tbl lineorder.tbl | sort | python part2Mapper.py | sort -n | python part2Reducer.py | python part2Mapper2.py | sort | python part2Reducer2.py

```
[ec2-user@ip-172-31-12-205 ~]$ cat customer.tbl lineorder.tbl | sort | python part2Mapper.py | sort -n | python part2Reducer.py | python part2Mapper2.py | sort | python part2Reducer2.py
ARGENTINA      2391736065660737
BRAZIL         2178049245933703
CANADA         2426493295203867
PERU           2316111574435684
UNITED STATES  2497485963211504
[ec2-user@ip-172-31-12-205 ~]$
```

However they fail when running first pass hadoop jar command. If you see why, please let me know.

PASS 1 command

```
hadoop jar /home/ec2-user/hadoop-2.6.4/hadoop-streaming-2.6.4.jar -D
mapred.text.key.comparator.options=-n -input /data/joinCustomer -output /data/outputCustomer007 -
mapper part2Mapper.py -file /home/ec2-user/part2Mapper.py -reducer part2Reducer.py -file /home/ec2-
user/part2Reducer.py
```

```

Total committed heap usage (bytes)=575668224
File Input Format Counters
    Bytes Read=194500959
19/06/17 02:24:09 ERROR streaming.StreamJob: Job not successful!
Streaming Command Failed!
[ec2-user@ip-172-31-12-205 ~]$
```

PASS 2 command would be:

```
hadoop jar /home/ec2-user/hadoop-2.6.4/hadoop-streaming-2.6.4.jar -input
/data/outputCustomer007/part-00000 -output /data/outputCustomerPass2 -mapper part2Mapper2.py -file
/home/ec2-user/part2Mapper2.py -reducer part2Reducer2.py -file /home/ec2-user/part2Reducer2.py
```

Testing on smaller samples

hadoop fs -mkdir -p /data/joinCustomerSample

```
[ec2-user@ip-172-31-12-205 ~]$ hadoop fs -ls /data/ | grep join
drwxr-xr-x  - ec2-user supergroup          0 2019-06-15 03:49 /data/joinCustomer
drwxr-xr-x  - ec2-user supergroup          0 2019-06-17 02:29 /data/joinCustomerSample
```

hadoop fs -put customer.tbl.sample lineorder.tbl.sample /data/joinCustomerSample

```
[ec2-user@ip-172-31-12-205 ~]$ hadoop fs -put customer.tbl.sample lineorder.tbl.sample /data/joinCustomerSample
[ec2-user@ip-172-31-12-205 ~]$ hadoop fs -ls /data/joinCustomerSample
Found 2 items
-rw-r--r--  2 ec2-user supergroup          46311 2019-06-17 02:31 /data/joinCustomerSample/customer.tbl.sample
-rw-r--r--  2 ec2-user supergroup          47530 2019-06-17 02:31 /data/joinCustomerSample/lineorder.tbl.sample
[ec2-user@ip-172-31-12-205 ~]$
```

**hadoop jar /home/ec2-user/hadoop-2.6.4/hadoop-streaming-2.6.4.jar -D
mapred.text.key.comparator.options=-n -input /data/joinCustomerSample -output
/data/outputCustomerSample002 -mapper part2Mapper.py -file /home/ec2-user/part2Mapper.py -reducer
par2Reducer.py -file /home/ec2-user/part2Reducer.py**

```
Virtual memory (bytes) snapshot=4286844928
Total committed heap usage (bytes)=404226048
File Input Format Counters
Bytes Read=93841
19/06/17 02:40:09 ERROR streaming.StreamJob: Job not successful!
Streaming Command Failed!
[ec2-user@ip-172-31-12-205 hadoop-2.6.4]$
```

Pig

```
cd $PIG_HOME
```

```
bin/pig
```

```
hadoop fs -put customer.tbl /user/ec2-user/
```

```
hadoop fs -put lineorder.tbl /user/ec2-user/
```

```
[ec2-user@ip-172-31-12-205 ~]$ hadoop fs -ls lineorder.tbl /user/ec2-user/ | grep .tbl
-rw-r--r--  2 ec2-user supergroup  594313001 2019-06-13 03:28 lineorder.tbl
-rw-r--r--  2 ec2-user supergroup   2837046 2019-06-13 03:28 /user/ec2-user/customer.tbl
-rw-r--r--  2 ec2-user supergroup  594313001 2019-06-13 03:28 /user/ec2-user/lineorder.tb
-rw-r--r--  2 ec2-user supergroup   17139259 2019-06-11 02:36 /user/ec2-user/part.tbl
drwxr-xr-x  - ec2-user supergroup           0 2019-06-12 04:09 /user/ec2-user/partdataswitc
hed.tbl
[ec2-user@ip-172-31-12-205 ~]$
```

```
customerData = LOAD '/user/ec2-user/customer.tbl' USING PigStorage('|')
AS (c_custkey:INT, c_name:CHARARRAY, c_address:CHARARRAY, c_city:CHARARRAY, c_nation:CHARARRAY,
c_region:CHARARRAY, c_phone:CHARARRAY, c_mktsegment:CHARARRAY);
```

```
DESCRIBE customerData;
```

```
grunt> DESCRIBE customerData;
customerData: {c_custkey: int,c_name: chararray,c_address: chararray,c_city: chararray,c_na
tion: chararray,c_region: chararray,c_phone: chararray,c_mktsegment: chararray}
grunt>
```

```
lineorderData = LOAD '/user/ec2-user/lineorder.tbl' USING PigStorage('|')
AS (lo_orderkey:INT, lo_linenummer:INT, lo_custkey:INT, lo_partkey:INT, lo_suppkey:INT, lo_orderdate:INT,
lo_orderpriority:CHARARRAY, lo_shippriority:CHARARRAY, lo_quantity:INT, lo_extendedprice:INT,
lo_ordertotalprice:INT, lo_discount:INT, lo_revenue:INT, lo_supplycost:INT, lo_tax:INT, lo_commitdate:INT,
lo_shipmode: CHARARRAY);
```

```
DESCRIBE lineorderData;
```

```
grunt> DESCRIBE lineorderData;
lineorderData: {lo_orderkey: int,lo_linenummer: int,lo_custkey: int,lo_partkey: int,lo_supp
key: int,lo_orderdate: int,lo_orderpriority: chararray,lo_shippriority: chararray,lo quanti
ty: int,lo_extendedprice: int,lo_ordertotalprice: int,lo_discount: int,lo_revenue: int,lo_s
upplycost: int,lo_tax: int,lo_commitdate: int,lo_shipmode: chararray}
grunt>
```

```

filter_lo = FILTER lineorderData BY lo_discount >= 4 AND lo_discount <= 6;
filter_cust = FILTER customerData BY c_region == 'AMERICA';
join_key = JOIN filter_lo BY (lo_custkey), filter_cust BY (c_custkey);
group_by = GROUP join_key BY (c_nation);
out = FOREACH group_by GENERATE group, SUM(join_key.lo_revenue) as revenue;
DESCRIBE out;

```

```

grunt> filter_lo = FILTER lineorderData BY lo_discount >= 4 AND lo_discount <= 6;
grunt> filter_cust = FILTER customerData BY c_region == 'AMERICA';
grunt> join_key = JOIN filter_lo BY (lo_custkey), filter_cust BY (c_custkey);
grunt> group_by = GROUP join_key BY (c_nation);
grunt> out = FOREACH group_by GENERATE group, SUM(join_key.lo_revenue) as revenue;
grunt> DESCRIBE out;
out: {group: chararray,revenue: long}
grunt>

```

DUMP out;

```

RedUtil - Total input paths to process : 1
(PERU,228441124985)
(BRAZIL,225595365795)
(CANADA,240715548308)
(ARGENTINA,243988697072)
(UNITED STATES,244263170830)
grunt>

```

STORE out INTO 'Part2Pig' USING PigStorage(',');

```

2019-06-13 03:38:06,604 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Success!
grunt>

```

hadoop fs -cat /user/ec2-user/Part2Pig/part-r-00000

```

[ec2-user@ip-172-31-12-205 ~]$ hadoop fs -cat /user/ec2-user/Part2Pig/part-r-00000
PERU,228441124985
BRAZIL,225595365795
CANADA,240715548308
ARGENTINA,243988697072
UNITED STATES,244263170830
[ec2-user@ip-172-31-12-205 ~]$

```

Part 3: Clustering (30 pts)

Create a new numeric file with 125,000 rows and 3 columns, separated by space – you can generate numeric data as you prefer, but submit the code that you have used.

```
ec2-user@ip-172-31-12-205:~/apache-mahout-distribution-0.11.2
GNU nano 2.9.8 numericfile.py

import random

for x in range(125000):
    val1 = str(random.randint(1,100))
    val2 = str(random.randint(1,100))
    val3 = str(random.randint(1,100))
    out = val1 + " " + val2 + " " + val3

    print out
```

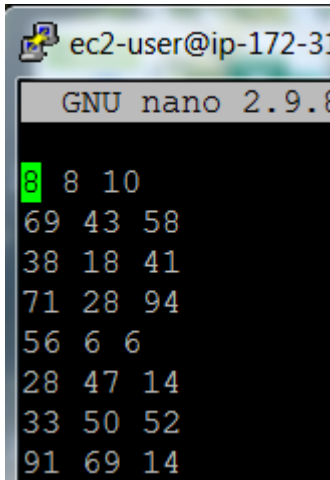
```
[ec2-user@ip-172-31-12-205 apache-mahout-distribution-0.11.2]$ python numericfile.py > numericfile.txt
```

- A. (5 pts) Using Mahout synthetic clustering as you have in a previous assignment on sample data. This entails running the same clustering command, but substituting your own input data and the right number of clusters.

`hadoop fs -mkdir -p testdata`

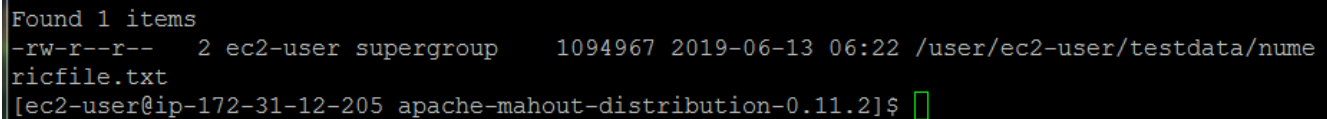
```
drwxr-xr-x - ec2-user supergroup 0 2019-06-13 06:11 /user/ec2-user/testdata
[ec2-user@ip-172-31-12-205 apache-mahout-distribution-0.11.2]$
```

Nano numericfile.txt



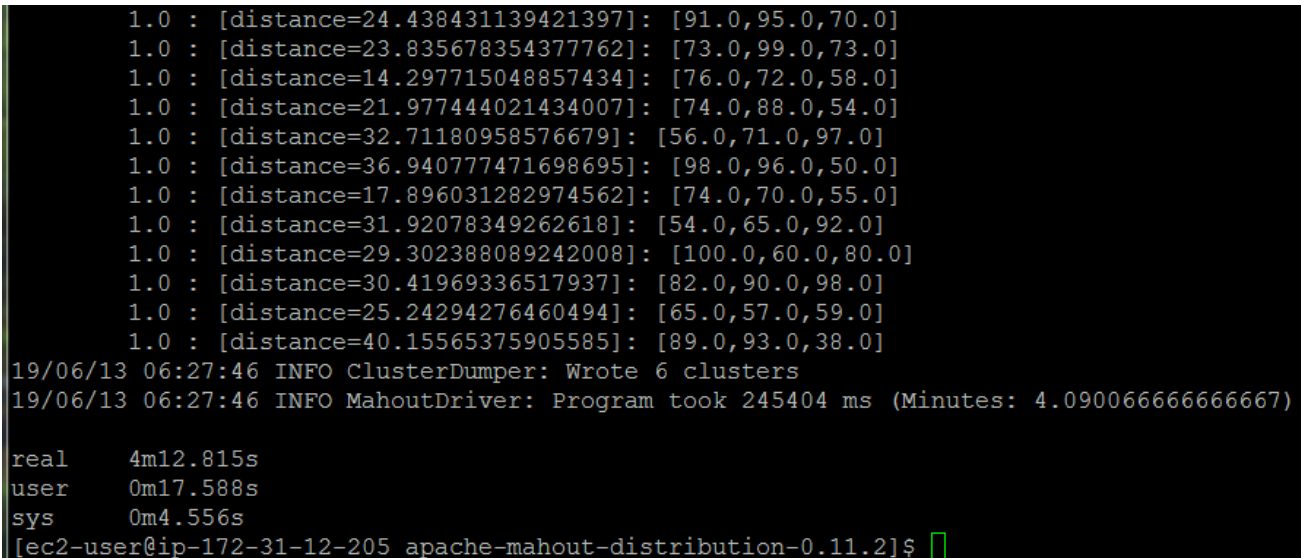
```
ec2-user@ip-172-31-12-205
GNU nano 2.9.8
8 8 10
69 43 58
38 18 41
71 28 94
56 6 6
28 47 14
33 50 52
91 69 14
```

hadoop fs -put numericfile.txt testdata/



```
Found 1 items
-rw-r--r--  2 ec2-user supergroup    1094967 2019-06-13 06:22 /user/ec2-user/testdata/numericfile.txt
[ec2-user@ip-172-31-12-205 apache-mahout-distribution-0.11.2]$
```

time mahout org.apache.mahout.clustering.syntheticcontrol.kmeans.Job



```
1.0 : [distance=24.438431139421397]: [91.0,95.0,70.0]
1.0 : [distance=23.835678354377762]: [73.0,99.0,73.0]
1.0 : [distance=14.297715048857434]: [76.0,72.0,58.0]
1.0 : [distance=21.977444021434007]: [74.0,88.0,54.0]
1.0 : [distance=32.71180958576679]: [56.0,71.0,97.0]
1.0 : [distance=36.940777471698695]: [98.0,96.0,50.0]
1.0 : [distance=17.896031282974562]: [74.0,70.0,55.0]
1.0 : [distance=31.92078349262618]: [54.0,65.0,92.0]
1.0 : [distance=29.302388089242008]: [100.0,60.0,80.0]
1.0 : [distance=30.41969336517937]: [82.0,90.0,98.0]
1.0 : [distance=25.24294276460494]: [65.0,57.0,59.0]
1.0 : [distance=40.15565375905585]: [89.0,93.0,38.0]
19/06/13 06:27:46 INFO ClusterDumper: Wrote 6 clusters
19/06/13 06:27:46 INFO MahoutDriver: Program took 245404 ms (Minutes: 4.090066666666667)

real    4m12.815s
user    0m17.588s
sys     0m4.556s
[ec2-user@ip-172-31-12-205 apache-mahout-distribution-0.11.2]$
```

```
mahout clusterdump --input output/clusters-10-final --pointsDir output/clusteredPoints --output clusteranalyze.txt
```

```
19/06/13 06:30:12 INFO AbstractJob: Command line arguments: {--dictionaryType=[text], --distanceMeasure=[org.apache.mahout.common.distance.SquaredEuclideanDistanceMeasure], --endPhase=[2147483647], --input=[output/clusters-10-final], --output=[clusteranalyze.txt], --outputFormat=[TEXT], --pointsDir=[output/clusteredPoints], --startPhase=[0], --tempDir=[temp]}
19/06/13 06:30:15 INFO ClusterDumper: Wrote 6 clusters
19/06/13 06:30:15 INFO MahoutDriver: Program took 2476 ms (Minutes: 0.04126666666666667)
[ec2-user@ip-172-31-12-205 apache-mahout-distribution-0.11.2]$
```

hadoop fs -ls output

```
[ec2-user@ip-172-31-12-205 apache-mahout-distribution-0.11.2]$ hadoop fs -ls output
Found 15 items
-rw-r--r--  2 ec2-user supergroup      194 2019-06-13 06:27 output/_policy
drwxr-xr-x  - ec2-user supergroup      0 2019-06-13 06:27 output/clusteredPoints
drwxr-xr-x  - ec2-user supergroup      0 2019-06-13 06:23 output/clusters-0
drwxr-xr-x  - ec2-user supergroup      0 2019-06-13 06:24 output/clusters-1
drwxr-xr-x  - ec2-user supergroup      0 2019-06-13 06:27 output/clusters-10-final
drwxr-xr-x  - ec2-user supergroup      0 2019-06-13 06:24 output/clusters-2
drwxr-xr-x  - ec2-user supergroup      0 2019-06-13 06:24 output/clusters-3
drwxr-xr-x  - ec2-user supergroup      0 2019-06-13 06:25 output/clusters-4
drwxr-xr-x  - ec2-user supergroup      0 2019-06-13 06:25 output/clusters-5
drwxr-xr-x  - ec2-user supergroup      0 2019-06-13 06:26 output/clusters-6
drwxr-xr-x  - ec2-user supergroup      0 2019-06-13 06:26 output/clusters-7
drwxr-xr-x  - ec2-user supergroup      0 2019-06-13 06:26 output/clusters-8
drwxr-xr-x  - ec2-user supergroup      0 2019-06-13 06:27 output/clusters-9
drwxr-xr-x  - ec2-user supergroup      0 2019-06-13 06:23 output/data
drwxr-xr-x  - ec2-user supergroup      0 2019-06-13 06:23 output/random-seeds
[ec2-user@ip-172-31-12-205 apache-mahout-distribution-0.11.2]$
```

- B. (25 pts) Using Hadoop streaming perform two iterations manually using 7 centers (initially with randomly chosen centers). As discussed in class, this would require passing a text file with cluster centers using `-file` option, opening the centers.txt in the mapper with `open('centers.txt', 'r')` and assigning a key to each point based on which center is the closest to each particular point. Your reducer would then compute the new centers, and at that point the iteration is done and the output of the reducer with new centers can be given to the next pass of the same code.

The only difference between first and second iteration is that in first iteration you have to pick the initial centers. In the 2nd iteration, the centers will be given to you by a previous pass of KMeans.

```
cp /home/ec2-user/centers.txt ./part3centers.txt
```

```
hadoop fs -put part3numericfile.txt /data/
```

```
ec2-user@ip-172-31-12-205:~  
GNU nano 2.9.8  
1 2 93  
54 31 5  
14 9 60  
39 30 6  
5 49 63  
88 5 68  
4 61 95
```

Mapper: part3Mapper.py

```
ec2-user@ip-172-31-12-205:~  
GNU nano 2.9.8 part3Mapper.py  
#!/usr/bin/python  
  
import os  
import sys  
  
#read initial centers from file  
fd = open('part3centers.txt', 'r')  
centers = []  
  
for line in fd:  
    l = line.strip()  
    v = line.split(' ')  
    centers.extend([v])  
fd.close()  
  
# read numericfile.txt  
for line in sys.stdin:  
    line = line.strip()  
    vals = line.split(' ')  
    v0 = float(vals[0])  
    v1 = float(vals[1])  
    v2 = float(vals[2])
```



```
cennum = None
calcdist = None
i = 0
for center in centers:
    c0 = float(center[0])
    c1 = float(center[1])
    c2 = float(center[2])
    dist = (v0-c0)**2 + (v1-c1)**2 + (v2-c2)**2 # distance
    #print str(dist)
    ed = dist**(0.5) # calculate square root, euclidian distance
    #print str(ed) # works ok
    if cennum:
        if ed < calcdist:
            cennum = i + 1
            calcdist = ed
    else:
        cennum = i + 1
        calcdist = ed
    i+= 1

print cennum, '\t', v0, '\t', v1, '\t', v2
```

Reducer: part3Reducer.py

```
ec2-user@ip-172-31-12-205:~  
GNU nano 2.9.8 part3Reducer.py  
#!/usr/bin/python  
import sys  
import os  
  
curr_id = None  
curr_cnt = 0  
id = None  
dim1 = []  
dim2 = []  
dim3 = []  
  
#delete part3centers.txt so par3Mapper can get new centers  
#os.remove("/home/ec2-user/part3centers.txt")  
#print("File Removed!")  
  
# The input comes from standard input (line by line)  
for line in sys.stdin:  
    line = line.strip()  
    #print 'line: ', line  
    # parse the line and split it by '\t'  
    ln = line.split('\t')  
    #print 'ln: ', ln  
    # grab the key (int)  
    id = int(ln[0].strip())  
    ln1= float(ln[1].strip())  
    ln2= float(ln[2].strip())  
    ln3= float(ln[3].strip())  
    #print 'id: ', id  
    #print 'ln1: ', ln1  
    #print 'ln2: ', ln2  
    #print 'ln3: ', ln3  
    #print 'curr_id: ', curr_id  
  
    if curr_id == id:  
        curr_cnt += 1  
        dim1.append(ln1) # list of values in vector first dimension  
        #print 'dim1: ', dim1  
        #print 'length of dim1:' , len(dim1)  
        dim2.append(ln2) # second dimension  
        dim3.append(ln3) # third dimension  
  
    else:  
        if curr_id: # output the count, single key completed  
            # NOTE: Change this to '%s\t%d' if your key is a string  
            #print 'len: ', dim1
```

```

        avrdim1 = str(sum(dim1)/len(dim1))
        avrdim2 = str(sum(dim2)/len(dim2))
        avrdim3 = str(sum(dim3)/len(dim3))
        out = avrdim1 + " " + avrdim2 + " " + avrdim3
        print out
        #print '%d\t%d' % (curr_id, curr_cnt)
        file = open("/home/ec2-user/part3centers.txt", "a")
        file.write(out + '\n')
        file.close()

    curr_id = id
    #curr_cnt = 0

# output the last key
if curr_id == id:
    avrdim1 = str(sum(dim1)/len(dim1))
    avrdim2 = str(sum(dim2)/len(dim2))
    avrdim3 = str(sum(dim3)/len(dim3))
    out = avrdim1 + " " + avrdim2 + " " + avrdim3
    print out
    file = open("/home/ec2-user/part3centers.txt", "a")
    file.write(out + '\n')
    file.close()

```

1st run

hadoop jar /home/ec2-user/hadoop-2.6.4/hadoop-streaming-2.6.4.jar -input /data/part3numericfile.txt -file part3centers.txt -output /data/part3kmeans07 -mapper part3Mapper.py -file /home/ec2-user/part3Mapper.py -reducer part3Reducer.py -file /home/ec2-user/part3Reducer.py

```

WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=1099063
File Output Format Counters
  Bytes Written=299
19/06/14 03:33:55 INFO streaming.StreamJob: Output directory: /data/part3kmeans07
[ec2-user@ip-172-31-12-205 ~]$

```

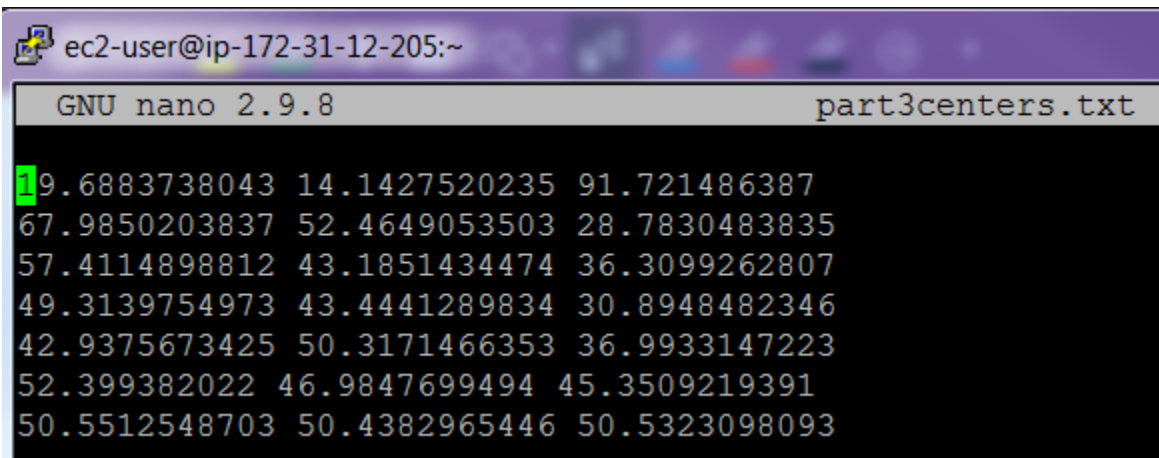
hadoop fs -cat /data/part3kmeans07/part-00000

```
[ec2-user@ip-172-31-12-205 ~]$ hadoop fs -cat /data/part3kmeans07/part-00000
19.6883738043 14.1427520235 91.721486387
67.9850203837 52.4649053503 28.7830483835
57.4114898812 43.1851434474 36.3099262807
49.3139754973 43.4441289834 30.8948482346
42.9375673425 50.3171466353 36.9933147223
52.399382022 46.9847699494 45.3509219391
50.5512548703 50.4382965446 50.5323098093
[ec2-user@ip-172-31-12-205 ~]$
```

rm part3centers.txt

hadoop fs -get /data/part3kmeans07/part-00000 part3centers.txt

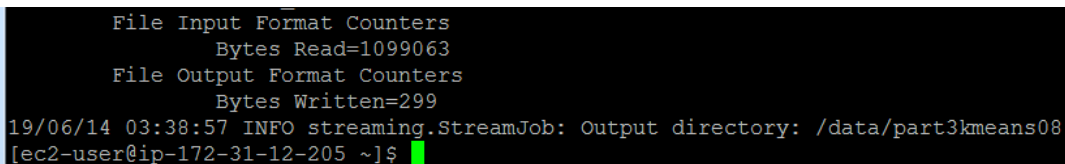
nano part3centers.txt



```
ec2-user@ip-172-31-12-205:~
GNU nano 2.9.8 part3centers.txt
19.6883738043 14.1427520235 91.721486387
67.9850203837 52.4649053503 28.7830483835
57.4114898812 43.1851434474 36.3099262807
49.3139754973 43.4441289834 30.8948482346
42.9375673425 50.3171466353 36.9933147223
52.399382022 46.9847699494 45.3509219391
50.5512548703 50.4382965446 50.5323098093
```

2nd run

hadoop jar /home/ec2-user/hadoop-2.6.4/hadoop-streaming-2.6.4.jar -input /data/part3numericfile.txt -file part3centers.txt -output /data/part3kmeans08 -mapper part3Mapper.py -file /home/ec2-user/part3Mapper.py -reducer part3Reducer.py -file /home/ec2-user/part3Reducer.py



```
File Input Format Counters
  Bytes Read=1099063
File Output Format Counters
  Bytes Written=299
19/06/14 03:38:57 INFO streaming.StreamJob: Output directory: /data/part3kmeans08
[ec2-user@ip-172-31-12-205 ~]$
```

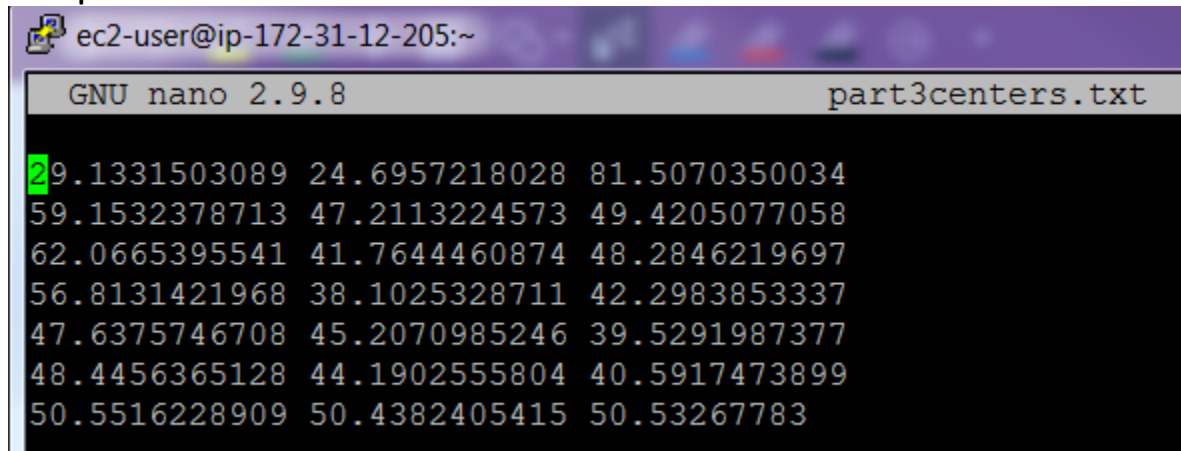
hadoop fs -cat /data/part3kmeans08/part-00000

```
19/06/14 03:38:57 INFO streaming.StreamJob: Output directory: /data/part3kmeans08
[ec2-user@ip-172-31-12-205 ~]$ hadoop fs -cat /data/part3kmeans08/part-00000
29.1331503089 24.6957218028 81.5070350034
59.1532378713 47.2113224573 49.4205077058
62.0665395541 41.7644460874 48.2846219697
56.8131421968 38.1025328711 42.2983853337
47.6375746708 45.2070985246 39.5291987377
48.4456365128 44.1902555804 40.5917473899
50.5516228909 50.4382405415 50.53267783
[ec2-user@ip-172-31-12-205 ~]$
```

rm part3centers.txt

hadoop fs -get /data/part3kmeans08/part-00000 part3centers.txt

nano part3centers.txt



```
ec2-user@ip-172-31-12-205:~
GNU nano 2.9.8 part3centers.txt
29.1331503089 24.6957218028 81.5070350034
59.1532378713 47.2113224573 49.4205077058
62.0665395541 41.7644460874 48.2846219697
56.8131421968 38.1025328711 42.2983853337
47.6375746708 45.2070985246 39.5291987377
48.4456365128 44.1902555804 40.5917473899
50.5516228909 50.4382405415 50.53267783
```

3rd run

hadoop jar /home/ec2-user/hadoop-2.6.4/hadoop-streaming-2.6.4.jar -input /data/part3numericfile.txt -file part3centers.txt -output /data/part3kmeans09 -mapper part3Mapper.py -file /home/ec2-user/part3Mapper.py -reducer part3Reducer.py -file /home/ec2-user/part3Reducer.py

```
WRONG_REDUCE=0
File Input Format Counters
    Bytes Read=1099063
File Output Format Counters
    Bytes Written=300
19/06/14 03:42:45 INFO streaming.StreamJob: Output directory: /data/part3kmeans09
[ec2-user@ip-172-31-12-205 ~]$
```

hadoop fs -cat /data/part3kmeans09/part-00000

```
19/06/14 03:42:45 INFO streaming.StreamJob: Output directory: /data/part3kmeans09
[ec2-user@ip-172-31-12-205 ~]$ hadoop fs -cat /data/part3kmeans09/part-00000
29.0024324778 31.0010484818 79.6794581446
51.8600925087 49.9528396357 69.5726479424
61.1794858603 44.9450875478 66.4815214968
62.3311082174 40.0939209102 58.0407374409
53.5155425631 44.2698349135 47.5726370799
53.143336317 43.9120722387 47.4426487513
50.5516388918 50.4384805549 50.5321498004
[ec2-user@ip-172-31-12-205 ~]$
```

rm part3centers.txt

hadoop fs -get /data/part3kmeans09/part-00000 part3centers.txt

nano part3centers.txt

```
ec2-user@ip-172-31-12-205:~
GNU nano 2.9.8 part3centers.txt
29.0024324778 31.0010484818 79.6794581446
51.8600925087 49.9528396357 69.5726479424
61.1794858603 44.9450875478 66.4815214968
62.3311082174 40.0939209102 58.0407374409
53.5155425631 44.2698349135 47.5726370799
53.143336317 43.9120722387 47.4426487513
50.5516388918 50.4384805549 50.5321498004
```

Submit a single document containing your written answers. Be sure that this document contains your name and "CSC 555 Project Phase 2" at the top.